

BY RYAN MURPHY

The AI-First CSM

How to build the kind of AI advantage inside your company that gets you cited by leadership — not just faster at your inbox.

"Ship the capability, and most people find their own way eventually. Ship the know-how alongside it, and adoption stops being a project."

RYAN MURPHY — CUSTOMER SUCCESS & AI
ENABLEMENT

10+ YEARS IN CS • 8+ IN SALES &
BD

A note from Ryan

I've spent ten years in customer success and eight before that in sales — most of it in SaaS, the last few years specifically in privacy and compliance software, which is about as far from "move fast and break things" as B2B software gets. I bring that up because the advice in this guide wasn't built in a lab or a hype cycle. It was built inside a job with real accounts, real renewals, and real people who had to actually use whatever I shipped.

This isn't a guide about prompting tricks. Plenty of people can already get a chatbot to draft a follow-up email. That's not an advantage anymore — it's table stakes. This guide is about something narrower and more useful: how to become the person in your company who **builds** the thing everyone else adopts, instead of the person who's just a little faster at their inbox.

Every idea in here has passed through one filter, over and over, for three years: **does this make someone's job easier, or did we just add another login?** If you remember nothing else, remember that question. It's the difference between AI as a feature you bolt on and AI as a genuine advantage you build.

What follows is the actual record — what I built, what it cost me to build it, and what it changed. Then the framework I use to find the next thing worth building. Then a straightforward process for doing this in your own company this week, whether you're a CSM, a BDR, or an AE. No theory you can't trace back to something that actually shipped.

Let's get into it.

— *Ryan*

Why AI-first, not AI-adjacent

Most people in customer success, sales, and BD are AI-adjacent. They use a chatbot to draft emails faster, summarize a call, or tidy up notes before a QBR. All useful. None of it changes how your company sees you — you're still doing the same job, just a little quicker at it.

AI-first is a different posture entirely. It's not about using the tool better than your peers — it's about being the person who **builds the reusable thing** that other people end up running on. A prompt library. A workflow. A small internal skill that flags risk or drafts outreach before anyone asks for it. You stop being a user of AI and start being the person who ships AI-shaped leverage for your whole team.

This distinction matters because of how careers actually move. Nobody gets promoted for answering emails faster — that's invisible by design; if it's working, no one notices. But people get cited by leadership, invited into strategy conversations, and asked to present in front of customers for the tool that quietly became "how we do things now." That's a fundamentally different kind of visibility, and it compounds in a way that faster typing never will.

So before anything tactical, the mindset shift: stop asking "how can AI help me do my job faster" and start asking a harder, more useful question.

"Does this make someone's job easier — or did we just add another login?"

That's the filter I run every idea through before I build it. It kills most ideas immediately, which is the point — the ones that survive it are the ones worth your time. The rest of this guide is the case study, and then the process, in that order. I want you to see what this actually looked like before I ask you to try it.

What AI-first actually built

None of this is theoretical. Here's what going AI-first looked like over three years, in a job most people wouldn't call a technical role.

STARTED AS A CHEAT SHEET

An eleven-section prompt library, now company-wide

I built myself a product-by-product reference for getting useful answers out of our platform's AI assistant — nothing official, just a document so I'd stop re-deriving the same prompts every week. It's now published in our internal Help Center, referenced in customer-facing webinars, and used across Customer Success, Sales, and Marketing.

The lesson: it never started as a mandate. It started as a personal irritant with a repeatable fix.

SALESFORCE- CONNECTED

An AI workflow that catches accounts before they go cold

A skill that scans a book of business for accounts that have gone quiet, flags them, and drafts the outreach — so the CSM's job becomes reviewing and sending, not staring at a dashboard wondering who to check on. It's used team-wide now, and leadership cites it as a concrete example of what AI does for the CS org, not just what it could do.

The lesson: the highest-leverage AI tools are the ones that replace a task nobody enjoyed in the first place.

~30 MINUTES, NOT WEEKS

A cross-region data migration that used to take weeks

Used AI-assisted tooling to move a customer's records between regions in a fraction of the usual time — with data quality that came in ahead of what a manual pass would have delivered.

The lesson: speed wasn't even the best part. The quality bar went up at the same time the timeline went down.

**30 DAYS,
START TO
FINISH****A full enterprise implementation, plus a migration off a competitor**

Stood up a new enterprise customer end-to-end on a 30-day timeline — on top of a separate, multi-month project moving 200-plus historical requests off a competitor's platform without losing continuity for the customer.

The lesson: AI-first doesn't mean skipping the hard, unglamorous migration work. It means having the tooling in place so the hard work doesn't also have to be slow.

**ONGOING
WEBINAR
SERIES****Becoming the voice customers hear, not just the team behind the scenes**

Co-presenting a recurring customer webinar series on getting real value from AI in privacy work — reporting, compliance gap analysis, risk assessment. The internal tooling didn't stay internal; it became the thing customers heard about directly from me.

The lesson: the internal build and the external reputation are the same project, just at different stages.

Look at that list again. None of it started with a leadership mandate to "go do AI." Every single item started as one person's irritation with a slow, repetitive, or error-prone process — solved quietly, then shown to someone, then adopted more widely because it actually worked. That path is repeatable. It's not luck, and it isn't specific to my company or my role. Section 4 is that path, broken into steps you can start this week.

But first — where do you even look for the next thing worth building? That's what the framework in the next section is for.

The four places I look for leverage

Everything I've built maps to one of four places I spend my attention. If you're not sure where to start looking, start here.

ONBOARDING

Time-to-value, not just go-live

Getting people live fast, without setting them up to stall out in month two. Go-live is a date on a calendar; time-to-value is whether the customer actually felt something get easier.

Look for: the manual, repetitive step in your onboarding checklist that AI could compress from a day to an hour.

ADOPTION

Turning early wins into renewals

Deepening usage across the product so the platform earns its renewal well before the renewal date does — instead of scrambling to justify the spend in the final quarter.

Look for: a way to surface silence before it becomes churn — the account that's gone quiet is the one nobody's watching.

ENABLEMENT

AI that people actually use

Building prompt libraries and workflows that lower the barrier between a powerful tool and the person who's supposed to benefit from it. Capability nobody uses isn't capability.

Look for: the tool your team already has access to but quietly avoids, and ask why.

PARTNERSHIP

Talking outcomes, not features

Connecting what you built to a metric leadership already tracks — retention, time-to-value, expansion — instead of a metric only you find interesting.

Look for: the dashboard your VP already checks every Monday. Build toward it, not around it.

You don't need all four at once. Pick the one closest to the work you already do — that's where you have the most credibility to build something and the shortest path to someone actually using it.

How to build your first AI asset

The prompt library that's now used company-wide at my company started as a personal cheat sheet. Here's the exact path from "my own irritation" to "the thing the team runs on," broken into steps.

- 1 Find the thing you redo more than twice a month**
It doesn't have to be hard. It has to be repetitive, slightly tedious, and something you're tired of doing manually. That's the tell.
- 2 Solve it for yourself first, quietly**
Don't ask permission to fix your own workflow. Build the prompt, the template, the small script. Use it for a few weeks before you show anyone.
- 3 Write down the how, not just the output**
This is the whole idea behind "ship the know-how alongside the capability." A result someone can't reproduce is a favor. A process someone can reuse is an asset.
- 4 Show it to one peer before you show your manager**
Find the person who has the exact same problem you had. If it helps them without you standing over their shoulder explaining it, it's ready for a wider audience. If it doesn't, you've just found your next fix.
- 5 Let usage make the case, not the announcement**
Don't pitch it to leadership as an idea. Let it already be in use by two or three people first. "Here's what we're already doing" is a completely different conversation than "here's what I think we should do."

The two questions I ask at the end of every session

This is the single habit I get asked about most, and it costs nothing to adopt. Before I close out any real working session with AI, I ask it two questions.

"What are you least confident about?"

"What's the biggest thing I don't realize about this situation?"

AI is confident by default — it will hand you a polished answer whether the underlying ground is solid or shaky, and it won't volunteer the difference unless you ask. These two questions force it to surface its own uncertainty and whatever it thinks you might be missing. That's usually exactly where the real risk — and, just as often, the real opportunity — is hiding.

I don't ask these idly. I ask them as a deliberate, closing step on anything that matters: a migration plan, a customer-facing deliverable, a workflow I'm about to hand off to someone else. It's saved me from shipping things that looked finished but weren't, more times than I can count.

Try it on the next real piece of work you do with AI — not a throwaway task, something with actual stakes. Ask both questions before you call it done. You'll be surprised how often the answer changes what you ship.

Turning a workflow hack into your AI brand

Building the thing is half the work. The other half is the quiet, unglamorous process of making sure it's the thing people associate with your name.

Document as you build, not after. A before/after screenshot, a rough note on time saved, a one-line description of the problem — capture it in the moment. You will not remember the details clearly enough six months later, when you actually need them for a review or a promotion case.

Demo it before it's polished. Offer to show it in a team meeting while it's still a little rough. "Here's something I've been building" gets people invested early — and gives you real feedback before you've sunk more time into the wrong version of it.

Find your internal audience before you find your executive sponsor. A tool that two peers are already using is a much easier pitch to leadership than an idea, however good, that no one has touched yet. Adoption is the credibility.

Say yes when the external ask comes. A webinar slot, a customer call, a conference panel — the moment your internal tool gets pulled outward is the moment it stops being "a thing I built at work" and starts being part of your professional reputation. That's the actual transition from AI-adjacent to AI-first, in career terms: your name and this capability become associated with each other, publicly, not just in your own head.

Where to start this week

Five questions. Answer them honestly and you'll have your first project by the time you're done.

1. The task I redo more than twice a month, that isn't hard but is tedious, is:

2. The person on my team who has the exact same problem is:

3. The one framework pillar this fits best — onboarding, adoption, enablement, or partnership — is:

4. The metric leadership already tracks that this connects to is:

5. Who I'll show this to first — not my manager — is:

Let's compare notes

If you're a CSM, a BDR, an AE, or a CS leader trying to get real traction with AI inside your company — not just a slide about it — I'd genuinely like to hear what's working for you. I'm especially glad to trade notes with people building their own prompt libraries, running AI-assisted migrations, or making the case for AI enablement to leadership for the first time.

Nothing in this guide is behind a paywall, and nothing about reaching out obligates you to anything. I'm not selling a course. I'm collecting the field notes of other people doing this for real, in their own companies, because that's how the next version of this guide gets better.

— *Ryan*

EMAIL rmurphy@trustarc.com

LINKEDIN [linkedin.com/in/ryan-murphy-52191629](https://www.linkedin.com/in/ryan-murphy-52191629)

Ship the capability. Then ship the know-how. That's the whole method.